

Neural and Evolutionary Computation (NEC)

Activity 1: Prediction with Back-Propagation and Linear Regression

Objective

Prediction using the following algorithms:

- **Neural Network with Back-Propagation (BP)**, implemented by the student.
- **Neural Network with Back-Propagation (BP-F)**, using free software.
- **Multiple Linear Regression (MLR-F)**, using free software.

Deliverables

This assignment can be done **alone or in pairs** (groups of two)

For this activity you must deliver **one PDF document** that includes:

- A link to the Github repository where the code of all the activity is accessible. More details on the code in the following sections.
- Explanations of the analysis of the data preprocessing, and the results of the executions that are detailed in the section below.
- The name of the file should be **A1-Name1Surname1-Name2Surname2.pdf**

Part 1: Selecting and analyzing the datasets

The predictions must be performed on three datasets:

1. File: *A1-turbine.txt*
 - 5 features: the first 4 are the input variables, the last one is the value to predict.
 - 451 patterns: use the first 85% for training and validation, and the remaining 15% for test.
2. File: *A1-synthetic.txt*
 - 10 features: the first 9 are the input variables, the last one is the value to predict.
 - 1000 patterns: use the first 80% for training and validation, and the remaining 20% for test.
3. Search a **dataset from the Internet**, with the following characteristics:
 - At least 6 features, one of them used for prediction.
 - The prediction variable must take real (float or double) values; it should not represent a categorical value (that would correspond to a classification task)
 - At least 400 patterns
 - Select randomly 80% of the patterns for training and validation, and the remaining 20% for test; it is important to shuffle the original data, to destroy any kind of sorting it could have.
 - In case of doubt on the suitability of the selected data, please ask the professor.

As an output of this part, **you should include in your report the following analysis:**

- For datasets 1 and 2, explain **what type of data normalization** you are going to apply to the input and output variables (the datasets are already cleaned, no need to preprocess them) and **why you apply it**.
- For dataset 3, **add to the documentation of this assignment the link to the source webpage**. Then explain what techniques of **data preprocessing** you are going to apply: check for missing values, represent correctly categorical values, and look for outliers. You should also consider if you must apply data normalization to some (or all) of the input / output variables.

Once this part is completed, you should **generate the preprocessed / normalized files** that are going to be the input of your analysis part.

Part 2: Implementation of BP

The most important part of this activity is the programming of a neural network with back propagation from scratch. Here are some instructions on how to code the project.

- You should **create a project in a github account** where you are going to develop the activity. You should do regular commits to this project so we can observe the evolution of the project from the initial file to the final delivery. We will apply a penalization to the grade if there is one single commit in the project.
- We recommend the implementation of the project using Python (version ≥ 3.6). We are providing a base code that can be used to start the project (MyNeuralNetwork.py).
- In this part you must implement all the methods necessary for the network to learn, you cannot use external libraries that already implement BP or Neural Networks.
- The **implementation** must be based on the algorithm and equations in **document [G] of Unit 3** at Moodle.
- The implementation must use the following **variables** to hold all the information about the structure of the multilayer neural network and the BP:
 - L : number of layers
 - n : an array with the number of units in each layer (including the input and output layers)
 - h : an array of arrays for the fields (h)
 - \mathbf{x}_i : an array of arrays for the activations (ξ)
 - w : an array of matrices for the weights (w)
 - θ : an array of arrays for the thresholds (θ)
 - Δ : an array of arrays for the propagation of errors (Δ)
 - d_w : an array of matrices for the changes of the weights (δw)
 - d_θ : an array of arrays for the changes of the weights ($\delta \theta$)
 - d_w_{prev} : an array of matrices for the previous changes of the weights, used for the momentum term ($\delta w^{(prev)}$)
 - d_θ_{prev} : an array of arrays for the previous changes of the thresholds, used for the momentum term ($\delta \theta^{(prev)}$)

- `fact`: the name of the activation function that it will be used. It can be one of these four: sigmoid, relu, linear, tanh.
- For example, the weight $w_{ij}^{(L)}$ between unit j in layer $L-1$ and unit i in layer L is accessed as `w[L][i,j]`
- The idea behind this structure is that the code must be able to deal with **arbitrary multilayer networks**. For example, a network with architecture 3:9:5:1 (4 layers, 3 input units, 1 output unit, and two hidden layers with 9 and 5 units, respectively), would have `n=[3; 9; 5; 1]`, and `xi` would be an array of length 4 (one component per layer), with `xi[1]` and array of real numbers of length 3, `xi[2]` and array of real numbers of length 9, `xi[3]` and array of real numbers of length 5, and `xi[4]` and array of real numbers of length 1. Similarly, `w[2]` would be an array 9×3 , `w[3]` an array 5×9 , and `w[4]` and array 1×5 ; `w[1]` is not used.
- Additionally, the use of this structure, name conventions and array dimensions make it easy to convert the equations into code.
- The code will receive one input dataset, and using the percentage of data that is passed as a parameter in the class constructor, should divide this dataset into training and validation. If the percentage is 0, then we consider that there is no validation, and all the input data is used for training.
- The class `MyNeuralNetwork` **will receive all these parameters in the class constructor**:
 - Number of layers
 - Number of units in each layer
 - Number of epochs
 - Learning rate and momentum
 - The selected activation function (sigmoid, relu, linear, tanh)
 - The percentage of data that should be used as the validation set
- The class `MyNeuralNetwork` **will provide three public functions** that can be called externally:
 - A function **`fit(X, y)`** that has 2 parameters: an array `X` of size (n_samples, n_features), which holds the training samples represented as floating point feature vectors; and a vector `y` of size (n_samples), which holds the target values (class labels) for the training samples. This method allows us to train the network with this data.
 - A function **`predict(X)`** that has 1 parameter, an array `X` of size (n_samples, n_features) that contains the samples. This method returns a vector with the predicted values for all the input samples.
 - A function **`loss_epochs()`** that returns 2 arrays of size (n_epochs, 2) that contain the evolution of the training error and the validation error for each of the epochs of the system, so this information can be plotted.

Part 3: Obtaining and comparing predictions using the three models (BP, BP-F, MLR-F)

This part is focused on comparing the results of the backpropagation model that has been developed against two open-source implementations of machine learning methods.

The coding corresponding on this part should be done using jupyter notebooks, which should also be included in the github of the project.

To study the quality of the predictions in this section, we are going to use two elements:

1. Although BP and MLR are based on the minimization of the mean squared error, it does not constitute a useful measure of the prediction error. Therefore, to compute the quality of a prediction, we will use the **mean absolute percentage error (MAPE)**, given by:

$$E(\%) = 100 \sum_{\mu} \left| \frac{y^{\mu} - z^{\mu}}{z^{\mu}} \right|$$

2. The best way to visualize the results is with **scatter plots** of the prediction value y^{μ} compared with the real value z^{μ} . The closer the points are to the diagonal, the better the prediction.

Part 3.1: Parameter comparison and selection

First, we need to choose the correct set of parameters that will provide the best prediction for the neural network that we have implemented. For this reason, for each of the three datasets we will have to explore what are the optimal values for the network parameters.

We have to explore some of the space of parameters, and the result of the prediction using those parameters. For this reason, you should include in the document **the following information for each of the 3 datasets**:

- A table that summarizes the quality of the prediction of each set of parameters using the MAPE value:

Number of layers	Layer Structure	Num epochs	Learning Rate	Momentum	Activation function	MAPE

- 2 or 3 representative scatter plots of the prediction value vs real value of some rows of this table (including the set of parameters that gives the minimum MAPE value).
- 2 or 3 plots of the evolution of the training and validation error as a function of the number of epochs (the information that can be obtained calling (the function `loss_epochs`)).
- A discussion/opinion of 1-2 paragraphs of why you think these parameters are the most adequate.

Part 3.2: Model result comparison

Second, we are going to compare the results obtained in our network against two already implemented models. We are going to compare the best results obtained in the previous section against the following models:

- A multi-linear regression from scikit-learn.
- A neural network model, which can be used from Tensorflow, Scikit-learn or any other python library.

For this part **you should include in the document**, for each of the three datasets, the following information:

- The description of the parameters used in each of the two additional models.
- A table comparing the prediction quality (MAPE) of the three models.
- The scatter plot of predicted vs real values for the three models.